

# Skip-gram Word2Vec

In this notebook, I'll lead you through using PyTorch to implement the [Word2Vec algorithm](https://en.wikipedia.org/wiki/Word2vec) (<https://en.wikipedia.org/wiki/Word2vec>) using the skip-gram architecture. By implementing this, you'll learn about embedding words for use in natural language processing. This will come in handy when dealing with things like machine translation.

## Readings

Here are the resources I used to build this notebook. I suggest reading these either beforehand or while you're working on this material.

- A really good [conceptual overview](http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model/) (<http://mccormickml.com/2016/04/19/word2vec-tutorial-the-skip-gram-model/>) of Word2Vec from Chris McCormick
- [First Word2Vec paper](https://arxiv.org/pdf/1301.3781.pdf) (<https://arxiv.org/pdf/1301.3781.pdf>) from Mikolov et al.
- [Neural Information Processing Systems, paper](http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf) (<http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf>) with improvements for Word2Vec also from Mikolov et al.

## Word embeddings

When you're dealing with words in text, you end up with tens of thousands of word classes to analyze; one for each word in a vocabulary. Trying to one-hot encode these words is massively inefficient because most values in a one-hot vector will be set to zero. So, the matrix multiplication that happens in between a one-hot input vector and a first, hidden layer will result in mostly zero-valued hidden outputs.

To solve this problem and greatly increase the efficiency of our networks, we use what are called **embeddings**. Embeddings are just a fully connected layer like you've seen before. We call this layer the embedding layer and the weights are embedding weights. We skip the multiplication into the embedding layer by instead directly grabbing the hidden layer values from the weight matrix. We can do this because the multiplication of a one-hot encoded vector with a matrix returns the row of the matrix corresponding the index of the "on" input unit.

$$\begin{bmatrix} 0 & 0 & 0 & 1 & 0 \end{bmatrix} \times \begin{bmatrix} 8 & 2 & 1 & 9 \\ 6 & 5 & 4 & 0 \\ 7 & 1 & 6 & 2 \\ 1 & 3 & 5 & 8 \\ 0 & 4 & 9 & 1 \end{bmatrix} = \begin{bmatrix} 1 & 3 & 5 & 8 \end{bmatrix}$$

Instead of doing the matrix multiplication, we use the weight matrix as a lookup table. We encode the words as integers, for example "heart" is encoded as 958, "mind" as 18094. Then to get hidden layer values for "heart", you just take the 958th row of the embedding matrix. This process is called an **embedding lookup** and the number of hidden units is the **embedding dimension**.

There is nothing magical going on here. The embedding lookup table is just a weight matrix. The embedding layer is just a hidden layer. The lookup is just a shortcut for the matrix multiplication. The lookup table is trained just like any weight matrix.

Embeddings aren't only used for words of course. You can use them for any model where you have a massive number of classes. A particular type of model called **Word2Vec** uses the embedding layer to find vector representations of words that contain semantic meaning.

## Word2Vec

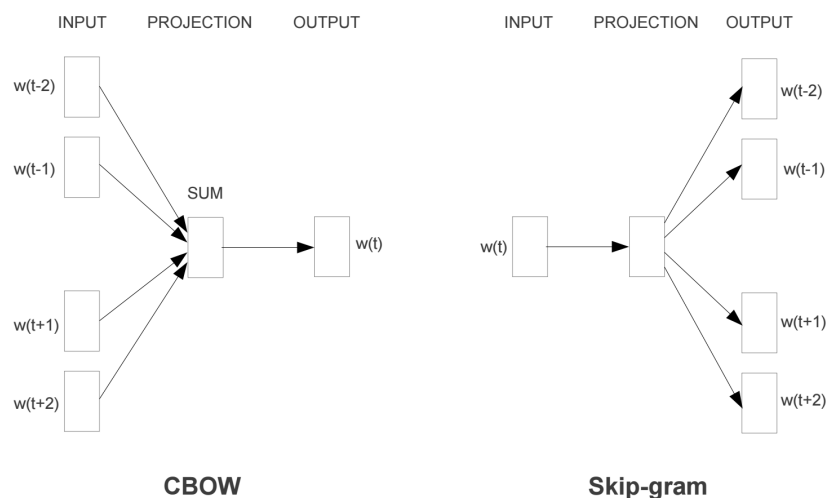
The Word2Vec algorithm finds much more efficient representations by finding vectors that represent the words. These vectors also contain semantic information about the words.

I often drink **coffee** in the mornings.  
 When I'm thirsty, I drink **water**.  
 I drink **tea**, before I go to sleep.

Words that show up in similar **contexts**, such as "coffee", "tea", and "water" will have vectors near each other. Different words will be further away from one another, and relationships can be represented by distance in vector space.

There are two architectures for implementing Word2Vec:

- CBOW (Continuous Bag-Of-Words) and
- Skip-gram



In this implementation, we'll be using the **skip-gram architecture** with **negative sampling** because it performs better than CBOW and trains faster with negative sampling. Here, we pass in a word and try to predict the words surrounding it in the text. In this way, we can train the network to learn representations for words that show up in similar contexts.

## Loading Data

Next, we'll ask you to load in data and place it in the `data` directory

1. Load the [text8 dataset \(https://s3.amazonaws.com/video.udacity-data.com/topher/2018/October/5bbe6499\\_text8/text8.zip\)](https://s3.amazonaws.com/video.udacity-data.com/topher/2018/October/5bbe6499_text8/text8.zip); a file of cleaned up *Wikipedia article text* from Matt Mahoney.
2. Place that data in the `data` folder in the home directory.
3. Then you can extract it and delete the archive, zip file to save storage space.

After following these steps, you should have one file in your data directory: `data/text8` .

In [1]:

```
# read in the extracted text file
with open('data/text8') as f:
    text = f.read()

# print out the first 100 characters
print(text[:100])
```

anarchism originated as a term of abuse first used against early work  
ing class radicals including t

## Pre-processing

Here I'm fixing up the text to make training easier. This comes from the `utils.py` file. The `preprocess` function does a few things:

- It converts any punctuation into tokens, so a period is changed to `<PERIOD>` . In this data set, there aren't any periods, but it will help in other NLP problems.
- It removes all words that show up five or *fewer* times in the dataset. This will greatly reduce issues due to noise in the data and improve the quality of the vector representations.
- It returns a list of words in the text.

This may take a few seconds to run, since our text file is quite large. If you want to write your own functions for this stuff, go for it!

In [2]:

```
import utils

# get list of words
words = utils.preprocess(text)
print(words[:30])
```

```
['anarchism', 'originated', 'as', 'a', 'term', 'of', 'abuse', 'first',
'used', 'against', 'early', 'working', 'class', 'radicals', 'includin
g', 'the', 'diggers', 'of', 'the', 'english', 'revolution', 'and', 'th
e', 'sans', 'culottes', 'of', 'the', 'french', 'revolution', 'whilst']
```

In [3]:

```
# print some stats about this word data
print("Total words in text: {}".format(len(words)))
print("Unique words: {}".format(len(set(words)))) # `set` removes any duplicate wor
```

Total words in text: 16680599  
Unique words: 63641

## Dictionaries

Next, I'm creating two dictionaries to convert words to integers and back again (integers to words). This is again done with a function in the `utils.py` file. `create_lookup_tables` takes in a list of words in a text and returns two dictionaries.

- The integers are assigned in descending frequency order, so the most frequent word ("the") is given the integer 0 and the next most frequent is 1, and so on.

Once we have our dictionaries, the words are converted to integers and stored in the list `int_words`.

In [4]:

```
vocab_to_int, int_to_vocab = utils.create_lookup_tables(words)
int_words = [vocab_to_int[word] for word in words]

print(int_words[:30])

[5233, 3080, 11, 5, 194, 1, 3133, 45, 58, 155, 127, 741, 476, 10571, 1
33, 0, 27349, 1, 0, 102, 854, 2, 0, 15067, 58112, 1, 0, 150, 854, 358
0]
```

## Subsampling

Words that show up often such as "the", "of", and "for" don't provide much context to the nearby words. If we discard some of them, we can remove some of the noise from our data and in return get faster training and better representations. This process is called subsampling by Mikolov. For each word  $w_i$  in the training set, we'll discard it with probability given by

$$P(w_i) = 1 - \sqrt{\frac{t}{f(w_i)}}$$

where  $t$  is a threshold parameter and  $f(w_i)$  is the frequency of word  $w_i$  in the total dataset.

Implement subsampling for the words in `int_words`. That is, go through `int_words` and discard each word given the probability  $P(w_i)$  shown above. Note that  $P(w_i)$  is the probability that a word is discarded. Assign the subsampled data to `train_words`.

In [5]:

```
from collections import Counter
import random
import numpy as np

threshold = 1e-5
word_counts = Counter(int_words)
#print(list(word_counts.items())[0]) # dictionary of int_words, how many times the

total_count = len(int_words)
freqs = {word: count/total_count for word, count in word_counts.items()}
p_drop = {word: 1 - np.sqrt(threshold/freqs[word]) for word in word_counts}
# discard some frequent words, according to the subsampling equation
# create a new list of words for training
train_words = [word for word in int_words if random.random() < (1 - p_drop[word])]

print(train_words[:30])

[127, 10571, 27349, 15067, 58112, 190, 10712, 104, 2731, 371, 2757, 68
6, 7088, 5233, 320, 44611, 5233, 2621, 8983, 4147, 141, 6437, 4186, 15
3, 5233, 1137, 6, 4860, 6753, 7573]
```

## Making batches

Now that our data is in good shape, we need to get it into the proper form to pass it into our network. With the skip-gram architecture, for each word in the text, we want to define a surrounding *context* and grab all the words in a window around that word, with size  $C$ .

From [Mikolov et al. \(https://arxiv.org/pdf/1301.3781.pdf\)](https://arxiv.org/pdf/1301.3781.pdf):

"Since the more distant words are usually less related to the current word than those close to it, we give less weight to the distant words by sampling less from those words in our training examples... If we choose  $C = 5$ , for each training word we will select randomly a number  $R$  in range  $[1 : C]$ , and then use  $R$  words from history and  $R$  words from the future of the current word as correct labels."

**Exercise:** Implement a function `get_target` that receives a list of words, an index, and a window size, then returns a list of words in the window around the index. Make sure to use the algorithm described above, where you chose a random number of words to from the window.

Say, we have an input and we're interested in the `idx=2` token, 741 :

```
[5233, 58, 741, 10571, 27349, 0, 15067, 58112, 3580, 58, 10712]
```

For  $R=2$  , `get_target` should return a list of four values:

```
[5233, 58, 10571, 27349]
```

In [6]:

```
def get_target(words, idx, window_size=5):
    ''' Get a list of words in a window around an index. '''

    R = np.random.randint(1, window_size+1)
    start = idx - R if (idx - R) > 0 else 0
    stop = idx + R
    target_words = words[start:idx] + words[idx+1:stop+1]

    return list(target_words)
```

In [7]:

```
# test your code!

# run this cell multiple times to check for random window selection
int_text = [i for i in range(10)]
print('Input: ', int_text)
idx=5 # word index of interest

target = get_target(int_text, idx=idx, window_size=5)
print('Target: ', target) # you should get some indices around the idx
```

```
Input:  [0, 1, 2, 3, 4, 5, 6, 7, 8, 9]
Target:  [3, 4, 6, 7]
```

## Generating Batches

Here's a generator function that returns batches of input and target data for our model, using the `get_target` function from above. The idea is that it grabs `batch_size` words from a words list. Then for each of those batches, it gets the target words in a window.

In [8]:

```
def get_batches(words, batch_size, window_size=5):
    ''' Create a generator of word batches as a tuple (inputs, targets) '''

    n_batches = len(words)//batch_size

    # only full batches
    words = words[:n_batches*batch_size]

    for idx in range(0, len(words), batch_size):
        x, y = [], []
        batch = words[idx:idx+batch_size]
        for ii in range(len(batch)):
            batch_x = batch[ii]
            batch_y = get_target(batch, ii, window_size)
            y.extend(batch_y)
            x.extend([batch_x]*len(batch_y))
        yield x, y
```

In [9]:

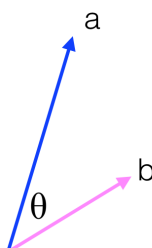
```
int_text = [i for i in range(20)]
x,y = next(get_batches(int_text, batch_size=4, window_size=5))

print('x\n', x)
print('y\n', y)
```

```
x
[0, 1, 1, 1, 2, 2, 2, 3, 3, 3]
y
[1, 0, 2, 3, 0, 1, 3, 0, 1, 2]
```

## Validation

Here, I'm creating a function that will help us observe our model as it learns. We're going to choose a few common words and few uncommon words. Then, we'll print out the closest words to them using the cosine similarity:



$$\text{similarity} = \cos(\theta) = \frac{\vec{a} \cdot \vec{b}}{|\vec{a}| |\vec{b}|}$$

We can encode the validation words as vectors  $\vec{a}$  using the embedding table, then calculate the similarity with each word vector  $\vec{b}$  in the embedding table. With the similarities, we can print out the validation words and words in our embedding table semantically similar to those words. It's a nice way to check that our embedding table is grouping together words with similar semantic meanings.

In [10]:

```
def cosine_similarity(embedding, valid_size=16, valid_window=100, device='cpu'):
    """ Returns the cosine similarity of validation words with words in the embedding table.
        Here, embedding should be a PyTorch embedding module.
    """

    # Here we're calculating the cosine similarity between some random words and
    # our embedding vectors. With the similarities, we can look at what words are
    # close to our random words.

    # sim = (a . b) / |a||b|

    embed_vectors = embedding.weight

    # magnitude of embedding vectors, |b|
    magnitudes = embed_vectors.pow(2).sum(dim=1).sqrt().unsqueeze(0)

    # pick N words from our ranges (0,window) and (1000,1000+window). lower id implies higher frequency
    valid_examples = np.array(random.sample(range(valid_window), valid_size//2))
    valid_examples = np.append(valid_examples,
                               random.sample(range(1000,1000+valid_window), valid_size//2))
    valid_examples = torch.LongTensor(valid_examples).to(device)

    valid_vectors = embedding(valid_examples)
    similarities = torch.mm(valid_vectors, embed_vectors.t())/magnitudes

    return valid_examples, similarities
```

## SkipGram model

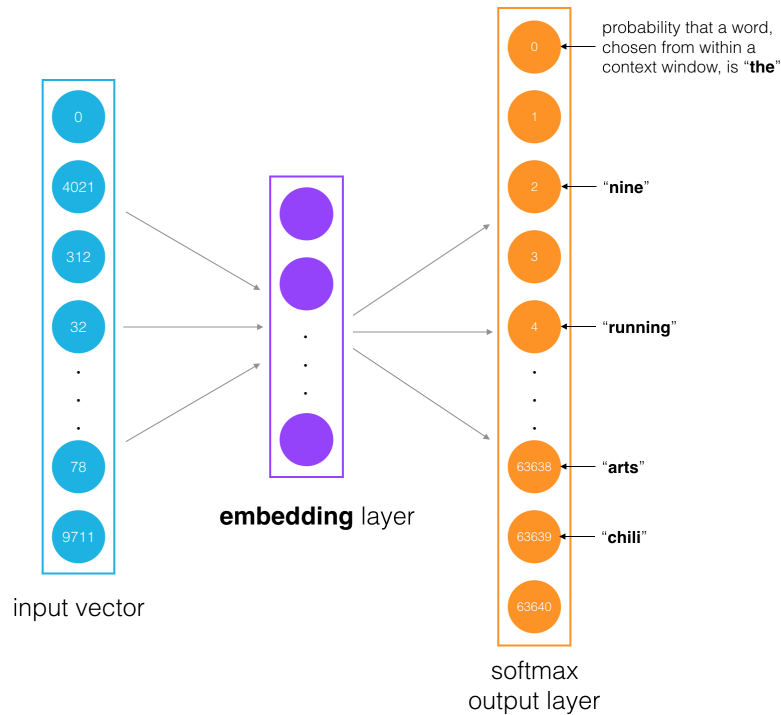
Define and train the SkipGram model.

You'll need to define an [embedding layer \(https://pytorch.org/docs/stable/nn.html#embedding\)](https://pytorch.org/docs/stable/nn.html#embedding) and a final, softmax output layer.

An Embedding layer takes in a number of inputs, importantly:

- **num\_embeddings** – the size of the dictionary of embeddings, or how many rows you'll want in the embedding weight matrix
- **embedding\_dim** – the size of each embedding vector; the embedding dimension

Below is an approximate diagram of the general structure of our network.



- The input words are passed in as batches of input word tokens.
- This will go into a hidden layer of linear units (our embedding layer).
- Then, finally into a softmax output layer.

We'll use the softmax layer to make a prediction about the context words by sampling, as usual.

## Negative Sampling

For every example we give the network, we train it using the output from the softmax layer. That means for each input, we're making very small changes to millions of weights even though we only have one true example. This makes training the network very inefficient. We can approximate the loss from the softmax layer by only updating a small subset of all the weights at once. We'll update the weights for the correct example, but only a small number of incorrect, or noise, examples. This is called "[negative sampling](http://papers.nips.cc/paper/5021-distributed-representations-of-words-and-phrases-and-their-compositionality.pdf)".

There are two modifications we need to make. First, since we're not taking the softmax output over all the words, we're really only concerned with one output word at a time. Similar to how we use an embedding table to map the input word to the hidden layer, we can now use another embedding table to map the hidden layer to the output word. Now we have two embedding layers, one for input words and one for output words. Secondly, we use a modified loss function where we only care about the true example and a small subset of noise examples.

$$-\log \sigma(u_{w_o}^\top v_{w_I}) - \sum_i^N \mathbb{E}_{w_i \sim P_n(w)} \log \sigma(-u_{w_i}^\top v_{w_I})$$

This is a little complicated so I'll go through it bit by bit.  $u_{w_o}^\top$  is the embedding vector for our "output" target word (transposed, that's the  $^\top$  symbol) and  $v_{w_I}$  is the embedding vector for the "input" word. Then the first term



$$\log \sigma \left( u_{w_O}^\top v_{w_I} \right)$$

says we take the log-sigmoid of the inner product of the output word vector and the input word vector. Now the second term, let's first look at

$$\sum_i^N \mathbb{E}_{w_i \sim P_n(w)}$$

This means we're going to take a sum over words  $w_i$  drawn from a noise distribution  $w_i \sim P_n(w)$ . The noise distribution is basically our vocabulary of words that aren't in the context of our input word. In effect, we can randomly sample words from our vocabulary to get these words.  $P_n(w)$  is an arbitrary probability distribution though, which means we get to decide how to weight the words that we're sampling. This could be a uniform distribution, where we sample all words with equal probability. Or it could be according to the frequency that each word shows up in our text corpus, the unigram distribution  $U(w)$ . The authors found the best distribution to be  $U(w)^{3/4}$ , empirically.

Finally, in

$$\log \sigma \left( -u_{w_i}^\top v_{w_I} \right),$$

we take the log-sigmoid of the negated inner product of a noise vector with the input vector.

$$-\log \underbrace{\sigma \left( u_{w_O}^\top v_{w_I} \right)}_{\substack{\uparrow \text{correct target}}} - \sum_i^N \mathbb{E}_{w_i \sim P_n(w)} \log \underbrace{\sigma \left( -u_{w_i}^\top v_{w_I} \right)}_{\substack{\downarrow \text{noisy target}}}$$

To give you an intuition for what we're doing here, remember that the sigmoid function returns a probability between 0 and 1. The first term in the loss pushes the probability that our network will predict the correct word  $w_O$  towards 1. In the second term, since we are negating the sigmoid input, we're pushing the probabilities of the noise words towards 0.

In [11]:

```
import torch
from torch import nn
import torch.optim as optim
```

In [12]:

```

class SkipGramNeg(nn.Module):
    def __init__(self, n_vocab, n_embed, noise_dist=None):
        super().__init__()

        self.n_vocab = n_vocab
        self.n_embed = n_embed
        self.noise_dist = noise_dist

        # define embedding layers for input and output words
        self.in_embed = nn.Embedding(n_vocab, n_embed)
        self.out_embed = nn.Embedding(n_vocab, n_embed)

        # Initialize embedding tables with uniform distribution
        # I believe this helps with convergence
        self.in_embed.weight.data.uniform_(-1, 1)
        self.out_embed.weight.data.uniform_(-1, 1)

    def forward_input(self, input_words):
        input_vectors = self.in_embed(input_words)
        return input_vectors

    def forward_output(self, output_words):
        output_vectors = self.out_embed(output_words)
        return output_vectors

    def forward_noise(self, batch_size, n_samples):
        """ Generate noise vectors with shape (batch_size, n_samples, n_embed) """
        if self.noise_dist is None:
            # Sample words uniformly
            noise_dist = torch.ones(self.n_vocab)
        else:
            noise_dist = self.noise_dist

        # Sample words from our noise distribution
        noise_words = torch.multinomial(noise_dist,
                                         batch_size * n_samples,
                                         replacement=True)

        device = "cuda" if model.out_embed.weight.is_cuda else "cpu"
        noise_words = noise_words.to(device)

        noise_vectors = self.out_embed(noise_words).view(batch_size, n_samples, sel

        return noise_vectors

```

In [13]:

```
class NegativeSamplingLoss(nn.Module):
    def __init__(self):
        super().__init__()

    def forward(self, input_vectors, output_vectors, noise_vectors):

        batch_size, embed_size = input_vectors.shape

        # Input vectors should be a batch of column vectors
        input_vectors = input_vectors.view(batch_size, embed_size, 1)

        # Output vectors should be a batch of row vectors
        output_vectors = output_vectors.view(batch_size, 1, embed_size)

        # bmm = batch matrix multiplication
        # correct log-sigmoid loss
        out_loss = torch.bmm(output_vectors, input_vectors).sigmoid().log()
        out_loss = out_loss.squeeze()

        # incorrect log-sigmoid loss
        noise_loss = torch.bmm(noise_vectors.neg(), input_vectors).sigmoid().log()
        noise_loss = noise_loss.squeeze().sum(1) # sum the losses over the sample

        # negate and sum correct and noisy log-sigmoid losses
        # return average batch loss
        return -(out_loss + noise_loss).mean()
```

## Training

Below is our training loop, and I recommend that you train on GPU, if available.

In [14]:

```

device = 'cuda' if torch.cuda.is_available() else 'cpu'

# Get our noise distribution
# Using word frequencies calculated earlier in the notebook
word_freqs = np.array(sorted(freqs.values(), reverse=True))
unigram_dist = word_freqs/word_freqs.sum()
noise_dist = torch.from_numpy(unigram_dist**(0.75)/np.sum(unigram_dist**(0.75)))

# instantiating the model
embedding_dim = 300
model = SkipGramNeg(len(vocab_to_int), embedding_dim, noise_dist=noise_dist).to(device)

# using the loss that we defined
criterion = NegativeSamplingLoss()
optimizer = optim.Adam(model.parameters(), lr=0.003)

print_every = 1500
steps = 0
epochs = 5

# train for some number of epochs
for e in range(epochs):

    # get our input, target batches
    for input_words, target_words in get_batches(train_words, 512):
        steps += 1
        inputs, targets = torch.LongTensor(input_words), torch.LongTensor(target_words)
        inputs, targets = inputs.to(device), targets.to(device)

        # input, output, and noise vectors
        input_vectors = model.forward_input(inputs)
        output_vectors = model.forward_output(targets)
        noise_vectors = model.forward_noise(inputs.shape[0], 5)

        # negative sampling loss
        loss = criterion(input_vectors, output_vectors, noise_vectors)

        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

    # loss stats
    if steps % print_every == 0:
        print("Epoch: {}/{}".format(e+1, epochs))
        print("Loss: ", loss.item()) # avg batch loss at this point in training
        valid_examples, valid_similarities = cosine_similarity(model.in_embed,
_, closest_idxxs = valid_similarities.topk(6)

        valid_examples, closest_idxxs = valid_examples.to('cpu'), closest_idxxs.to('cpu')
        for ii, valid_idx in enumerate(valid_examples):
            closest_words = [int_to_vocab[idx.item()] for idx in closest_idxxs[ii]]
            print(int_to_vocab[valid_idx.item()] + " | " + ', '.join(closest_words))
        print("...\n")

```

Epoch: 1/5

Loss: 6.7322540283203125

to | the, earth, barium, protection, and

zero | the, and, in, of, one

been | pathogen, competing, restored, advertising, demanded

between | herman, junta, kenyatta, two, outsiders  
also | beowulf, creatures, recommendation, cooler, boundary  
it | the, that, confer, send, as  
a | make, of, harper, that, the  
all | studio, blubber, drawbacks, accidentally, autonomous  
quite | neighbours, once, casinos, stuttgart, monster  
instance | marcus, teeth, leon, peninsular, married  
report | in, festivals, member, backbone, strains  
pressure | khan, muni, bipolar, swimming, satellite  
know | unify, had, klara, shapeshifting, households  
stage | clark, mayor, yates, supplements, doubtful  
engineering | sacs, z, one, heaven, passes  
animals | concentration, flicks, silvertown, absorb, geologically  
...

Epoch: 1/5

Loss: 4.835074424743652

his | in, the, a, of, as  
often | procurator, of, that, networks, established  
had | one, nine, which, zero, in  
many | of, which, rocks, with, two  
see | of, from, as, with, a  
and | the, a, to, of, in  
between | two, is, junta, neoliberalism, three  
people | with, in, for, alex, some  
construction | dutch, universiti, to, testament, atmospheres  
award | waqf, era, highway, by, bonding  
ocean | malthus, ubiquitous, regimens, who, transiting  
engine | unproductive, latina, zoe, by, battlefield  
quite | casinos, neighbours, once, portable, monster  
pope | comparison, coup, cheek, narrowly, mph  
event | soils, cdots, contingent, charlton, endless  
lived | disadvantage, saltwater, underwater, overview, donkeys  
...

Epoch: 1/5

Loss: 4.341375350952148

when | to, for, by, an, they  
or | is, be, a, the, in  
only | by, in, to, be, as  
known | a, and, nine, were, one  
not | to, be, is, that, it  
about | to, and, is, of, his  
are | is, that, a, in, the  
in | the, and, of, a, to  
rise | fax, laws, hells, comedienne, colonize  
magazine | liberal, delirium, pivot, narnia, taboos  
ocean | malthus, zeppelins, size, accuses, ubiquitous  
dr | cleopatra, chojn, suburb, aviator, runways  
experience | thought, god, affluence, interdiction, an  
construction | dutch, pun, of, atmospheres, universiti  
instance | marcus, nevertheless, gash, conflict, festivals  
bible | babylonian, vibrates, fielded, cbt, guilbert  
...

Epoch: 1/5

Loss: 3.6369738578796387

states | united, in, government, seven, state  
his | he, the, in, during, with  
many | most, by, of, the, from  
two | four, five, seven, zero, one

about | of, the, and, in, zero  
 war | had, united, the, one, of  
 they | the, are, a, of, that  
 eight | nine, one, seven, four, five  
 police | elections, states, administration, economic, countries  
 recorded | patass, toothed, mats, reuptake, d  
 pressure | temperature, use, muni, where, can  
 road | towns, volvo, gauge, hen, shakespearian  
 http | www, software, selecting, links, large  
 egypt | west, israel, european, hobbit, afghanistan  
 alternative | wikis, imbedded, davenport, slowly, mingus  
 institute | chocolates, summations, incurring, brushed, christianity  
 ...

Epoch: 1/5

Loss: 3.504284381866455

is | a, the, such, to, of  
 after | and, nine, one, was, the  
 have | the, are, from, most, that  
 nine | one, eight, four, seven, six  
 may | of, with, their, the, and  
 time | a, many, in, by, most  
 see | of, by, are, and, is  
 with | a, the, in, of, an  
 creation | yet, god, world, since, thorns  
 cost | systems, rate, blackened, expensive, taint  
 something | knowledge, they, follow, question, comments  
 existence | human, our, argument, teachings, principle  
 smith | charles, american, william, richard, politician  
 governor | charles, president, general, queen, executive  
 police | united, courts, military, prisoners, u  
 joseph | writer, ii, french, born, d  
 ...

Epoch: 1/5

Loss: 3.3184025287628174

some | are, that, is, still, to  
 more | these, are, which, to, or  
 he | his, him, her, after, wife  
 other | are, and, or, for, is  
 th | century, history, west, five, roman  
 however | to, the, which, some, that  
 such | is, different, or, are, be  
 would | to, was, they, him, did  
 http | www, links, com, web, website  
 creation | pushcha, since, they, to, beginning  
 instance | can, not, functions, complex, value  
 something | that, you, your, could, anything  
 bible | hebrew, testament, christian, jewish, biblical  
 bill | american, film, nine, george, singer  
 file | format, user, files, software, computer  
 quite | have, different, are, such, can  
 ...

Epoch: 2/5

Loss: 2.8340201377868652

six | four, eight, seven, one, three  
 can | be, or, techniques, types, using  
 between | and, from, part, thus, the  
 have | some, but, the, in, are  
 war | military, forces, troops, killed, soviet

from | into, and, which, the, in  
 when | to, the, can, put, them  
 history | links, see, external, list, and  
 brother | wife, died, father, his, son  
 mathematics | mathematical, theory, theories, science, theorem  
 consists | are, or, note, and, terms  
 event | football, ever, park, saw, ball  
 experience | towards, effects, that, what, belief  
 police | campaign, convicted, military, armed, party  
 versions | classic, video, features, version, software  
 egypt | dynasty, persian, east, empire, egyptian  
 ...

Epoch: 2/5

Loss: 2.5524561405181885

were | with, became, heavy, was, in  
 system | systems, code, computer, programs, using  
 other | or, these, make, see, often  
 over | a, and, has, there, or  
 but | and, that, however, an, of  
 was | had, after, victory, during, defeat  
 states | united, military, u, state, armed  
 during | was, in, first, the, s  
 except | divided, or, and, some, is  
 writers | works, fiction, literary, poets, published  
 something | might, could, our, without, more  
 road | city, street, park, east, washington  
 governor | president, appointed, elected, cabinet, minister  
 engine | engines, speed, models, cylinder, car  
 san | university, california, los, college, santa  
 frac | x, equation, f, mathbf, cdot  
 ...

Epoch: 2/5

Loss: 2.7377986907958984

its | the, between, and, more, their  
 was | he, had, daughter, days, of  
 it | thus, be, can, into, a  
 into | to, it, a, the, by  
 new | university, city, york, history, college  
 up | and, the, can, will, move  
 or | non, to, are, a, these  
 seven | one, eight, nine, four, zero  
 engineering | engineers, systems, technology, electrical, electronics  
 marriage | her, married, divorce, life, he  
 pope | roman, papal, rome, leo, catholics  
 shown | complex, finite, number, axis, above  
 egypt | syria, persian, egyptian, cairo, sinai  
 magazine | publishing, books, bbc, interview, career  
 grand | des, paris, royal, north, st  
 channel | satellite, radio, cable, coverage, broadcast  
 ...

Epoch: 2/5

Loss: 2.6995224952697754

no | import, done, jargon, info, duplicate  
 be | are, any, time, for, these  
 four | five, three, two, six, nine  
 five | four, six, three, zero, two  
 their | and, the, themselves, to, or  
 system | systems, data, programming, code, implementation

the | in, of, by, and, to  
 known | the, from, a, north, in  
 square | kilometers, situated, located, near, county  
 versions | computer, macintosh, version, feature, microsoft  
 magazine | media, newspapers, fiction, bbc, publishing  
 woman | her, she, husband, mother, birth  
 shows | television, tv, genre, game, final  
 assembly | elected, elections, legislative, cabinet, president  
 bill | office, president, executive, bush, tony  
 troops | army, military, war, armies, soviet  
 ...

Epoch: 2/5

Loss: 2.580954074859619

to | the, and, when, a, as  
 th | nd, century, one, zero, four  
 has | are, other, or, as, include  
 with | in, of, a, the, also  
 zero | five, three, four, two, nine  
 however | in, their, to, have, was  
 if | we, will, must, so, be  
 more | are, most, as, or, and  
 dr | ed, four, story, bibliography, john  
 recorded | album, recording, song, label, records  
 engine | engines, cylinder, combustion, fuel, prototype  
 rise | revolution, influence, thought, attempts, and  
 award | awards, winners, best, awarded, fame  
 issue | opinion, court, rejected, amendment, supreme  
 pressure | heat, gas, effective, forcing, air  
 versions | version, pc, macintosh, microsoft, os  
 ...

Epoch: 2/5

Loss: 2.7661571502685547

a | which, and, to, the, or  
 it | to, is, for, an, of  
 for | and, with, it, a, which  
 years | year, birth, months, death, female  
 would | he, wanted, to, him, was  
 however | to, some, during, it, occur  
 five | one, six, two, eight, seven  
 there | are, which, except, typically, in  
 ice | winter, hockey, temperature, temperatures, lake  
 pre | main, modern, see, millennium, early  
 behind | left, told, hitting, driving, doors  
 existence | defined, follows, interpretation, quantum, notion  
 governor | appointed, prime, minister, john, elected  
 scale | scales, temperature, flat, major, instruments  
 account | potential, hypothesis, how, this, values  
 freedom | political, accused, presidential, rights, social  
 ...

Epoch: 3/5

Loss: 2.850525379180908

about | people, what, part, according, life  
 zero | two, five, four, seven, eight  
 more | are, than, a, its, the  
 i | we, you, him, me, daughter  
 up | down, their, to, the, from  
 to | a, the, for, in, this  
 d | b, laureate, politician, writer, actor



were | their, had, forced, took, to  
something | we, you, want, sufficiently, feels  
marriage | married, marriages, marry, divorce, wife  
lived | died, throne, sons, his, mother  
pressure | temperature, heat, liquid, melting, pressures  
question | questions, ought, belief, argued, ethical  
channel | channels, cable, television, tv, radio  
report | reports, review, news, pdf, terrorist  
mainly | largest, ethnic, large, many, among  
...

Epoch: 3/5

Loss: 2.4977896213531494

there | are, is, every, least, any  
which | a, is, it, to, be  
would | to, him, that, their, lose  
between | this, into, non, formed, characterized  
can | a, be, is, if, cannot  
with | the, to, or, a, some  
were | was, had, before, in, who  
who | his, married, he, friend, had  
liberal | liberals, party, liberalism, conservative, democracy  
engineering | technology, engineers, systems, university, education  
rise | massive, increasing, decline, wealth, depends  
prince | duke, son, princess, sir, elizabeth  
troops | army, forces, war, germans, armies  
san | francisco, diego, california, jose, santa  
dr | robert, lynn, eric, journal, musician  
discovered | discovery, found, observations, periodic, obtained  
...

Epoch: 3/5

Loss: 2.225531578063965

or | are, is, to, all, other  
to | be, the, it, or, able  
seven | eight, nine, one, four, three  
have | be, were, and, to, some  
s | his, one, two, eight, nine  
world | most, the, cultural, war, race  
called | is, use, are, an, used  
where | n, t, to, left, above  
gold | silver, copper, metals, tin, diamonds  
frac | x, equation, cdot, mathbf, cos  
engine | engines, cylinder, fuel, piston, powered  
writers | philosophers, poets, novelists, deaths, literature  
numerous | in, regarding, other, many, led  
report | news, reports, links, cia, org  
joseph | john, william, born, benjamin, james  
ice | rock, frozen, hockey, snow, crush  
...

Epoch: 3/5

Loss: 2.6488029956817627

b | d, k, r, n, y  
will | if, would, can, not, should  
from | the, of, in, and, with  
as | a, such, the, or, in  
known | and, the, of, was, roman  
he | his, him, himself, her, who  
an | which, the, it, on, is  
can | used, are, it, be, if

freedom | social, rights, policies, political, promoting  
 gold | silver, copper, metals, precious, timber  
 issue | political, accepted, rejected, commentators, membership  
 discovered | discovery, observations, experiments, astronomer, found  
 nobel | prize, laureate, physicist, chemist, politician  
 governor | secretary, president, appointed, appoints, senate  
 ice | rock, crust, dry, frozen, rocks  
 dr | starring, michael, ed, allen, his  
 ...

Epoch: 3/5

Loss: 2.487758159637451

between | there, eastern, in, or, as  
 years | zero, female, age, three, birth  
 their | they, and, however, because, only  
 may | are, usually, means, in, or  
 over | the, under, zero, and, in  
 known | th, by, the, also, in  
 to | and, the, of, in, was  
 at | in, was, england, near, he  
 troops | forces, army, invasion, attack, invaded  
 dr | nine, four, allen, poet, bibliography  
 numerous | also, around, most, cited, and  
 professional | football, fame, american, school, canadian  
 construction | built, constructed, buildings, architectural, building  
 scale | scales, tend, materials, technological, produce  
 brother | son, daughter, cousin, sister, wife  
 ice | hockey, rock, skating, crust, frozen  
 ...

Epoch: 3/5

Loss: 2.500605583190918

have | their, are, this, were, there  
 their | have, they, some, or, are  
 use | such, these, applications, possible, used  
 six | one, two, eight, five, zero  
 but | it, or, be, the, that  
 his | he, father, him, brother, son  
 years | year, death, five, female, age  
 they | their, to, however, so, would  
 active | organized, radical, passive, bahasa, member  
 mathematics | mathematical, mathematicians, euclid, sciences, math  
 mean | n, mu, argument, variable, sin  
 running | run, platform, address, runs, on  
 alternative | music, for, slang, see, used  
 channel | channels, cable, radio, broadcast, television  
 shown | in, both, called, here, is  
 ice | hockey, skating, glacier, frozen, winter  
 ...

Epoch: 4/5

Loss: 2.3018791675567627

more | are, but, only, is, very  
 its | the, from, of, is, has  
 first | was, the, for, one, had  
 which | is, it, in, the, also  
 this | it, that, in, most, to  
 th | century, nd, centuries, eight, period  
 where | in, when, or, to, long  
 three | two, one, four, zero, six  
 orthodox | catholic, church, christians, christianity, orthodoxy

award | awards, nomination, oscar, awarded, best  
bbc | weekly, august, listing, broadcasts, news  
accepted | believe, according, opinion, interpretation, questions  
recorded | album, song, songs, albums, date  
square | density, miles, length, mi, km  
troops | army, war, forces, armies, battle  
professional | football, team, baseball, teams, players  
...

Epoch: 4/5

Loss: 2.524277687072754

of | in, and, the, by, s  
united | states, british, department, canada, countries  
in | of, and, the, was, from  
first | one, the, s, was, by  
his | he, him, himself, who, father  
see | of, is, history, in, links  
state | california, university, college, colorado, district  
has | the, in, which, to, with  
arts | school, art, schools, academy, technology  
mainly | widely, mostly, region, among, recent  
dr | michael, allen, ed, writer, richard  
smith | john, jr, adam, joe, mormon  
egypt | egyptian, bc, egyptians, arab, syria  
articles | links, org, dictionary, journal, discusses  
quite | have, more, are, usually, very  
ice | rock, temperatures, frozen, winter, surface  
...

Epoch: 4/5

Loss: 2.364931583404541

state | absolute, of, energy, system, constant  
not | be, that, does, without, to  
they | their, those, for, but, are  
of | the, in, and, is, a  
was | his, had, s, in, after  
s | his, one, in, a, of  
people | who, native, ethnic, peoples, nation  
also | the, in, for, of, is  
shows | television, shown, wave, series, combination  
taking | following, result, last, year, the  
magazine | published, news, newspapers, interview, publications  
operations | operation, numbers, algebra, operational, numerical  
pre | historical, other, use, cultural, historic  
account | accounts, about, notes, that, value  
heavy | heavier, metal, smoke, guns, manufacture  
file | files, windows, format, unix, rom  
...

Epoch: 4/5

Loss: 2.3778340816497803

where | it, at, and, not, all  
if | we, must, then, let, given  
after | until, the, during, by, was  
s | of, one, by, and, the  
may | or, does, not, that, the  
however | their, this, would, made, its  
was | were, the, by, had, to  
many | most, while, such, of, several  
smith | william, john, friedman, jr, j  
issue | official, ties, issues, advocate, accepted

resources | information, directory, list, links, mining  
 applied | materials, context, introduction, study, fundamental  
 professional | sports, amateur, football, players, team  
 running | platform, operating, larry, run, os  
 test | tests, team, ground, unmanned, hazardous  
 bbc | itv, broadcast, television, broadcasts, aired  
 ...

Epoch: 4/5

Loss: 2.4204654693603516

is | are, the, or, a, an  
 seven | one, five, three, four, six  
 often | or, more, usually, sometimes, especially  
 with | and, a, of, in, for  
 when | would, before, to, was, he  
 during | in, was, s, and, the  
 the | in, a, to, of, and  
 are | or, is, be, called, other  
 universe | cosmology, bang, god, worlds, creator  
 hit | hits, hitting, run, album, inning  
 square | kilometers, adjacent, km, central, mi  
 institute | university, research, science, college, institutes  
 pressure | thrust, pressures, cooled, temperature, measures  
 rise | decline, early, end, among, fall  
 grand | prix, duke, prestigious, battle, ney  
 gold | silver, metals, copper, precious, platinum  
 ...

Epoch: 4/5

Loss: 2.5025265216827393

can | be, which, that, or, is  
 in | the, of, by, and, which  
 use | used, using, systems, these, uses  
 and | the, of, in, a, to  
 had | was, were, he, to, s  
 known | by, a, in, is, an  
 they | to, still, not, but, them  
 a | the, which, or, is, and  
 resources | information, source, management, natural, processing  
 engine | engines, combustion, piston, fuel, turbine  
 behind | doors, door, mask, race, car  
 defense | military, defence, nato, training, armed  
 mathematics | mathematical, mathematicians, algebra, geometry, euclid  
 shown | is, the, seen, of, then  
 prince | crown, princess, throne, regent, empress  
 something | wrong, let, you, say, indeed  
 ...

Epoch: 5/5

Loss: 2.1604182720184326

can | or, be, cannot, useful, if  
 no | that, to, as, not, for  
 a | with, the, as, is, by  
 six | one, eight, four, seven, two  
 state | democratic, of, states, government, president  
 will | if, when, it, to, must  
 between | of, the, are, which, within  
 new | nine, s, three, history, york  
 primarily | modern, such, used, include, and  
 paris | france, french, jean, versailles, fran  
 discovered | discovery, found, geologist, been, specimens

liberal | conservative, liberalism, liberals, conservatives, social  
shown | effects, treatment, these, so, cancer  
heavy | vehicles, armoured, weapon, personnel, primarily  
hold | belief, all, doctrine, or, clearly  
pre | ancient, existed, see, millennium, archaeological  
...

Epoch: 5/5

Loss: 2.2580490112304688

his | he, him, himself, who, father  
so | to, when, they, but, that  
system | systems, provide, types, tool, multi  
of | the, in, and, a, is  
is | a, which, are, the, of  
zero | two, four, five, nine, three  
war | troops, allied, army, forces, wwii  
from | and, the, of, in, as  
san | francisco, diego, california, santa, antonio  
ice | snow, temperature, rock, winter, rocks  
operating | unix, microsoft, user, multitasking, software  
operations | operation, intelligence, covert, terrorist, military  
assembly | elected, legislative, unicameral, president, elections  
question | argument, arguments, which, answers, questions  
pressure | liquid, temperature, boiling, heating, gas  
nobel | prize, laureate, physicist, chemist, recipient  
...

Epoch: 5/5

Loss: 2.550434112548828

no | not, any, have, that, does  
eight | seven, five, nine, one, zero  
who | his, he, and, had, whom  
been | it, in, found, for, most  
at | the, in, from, five, four  
over | zero, one, the, three, four  
this | it, the, a, is, which  
on | in, a, the, s, for  
report | news, pdf, agency, review, cia  
shown | the, a, is, least, this  
units | unit, measured, si, density, measures  
recorded | recording, songs, records, period, year  
institute | university, college, universities, polytechnic, arts  
stage | movies, played, career, featured, hip  
lived | mother, life, briefly, she, and  
scale | scales, large, relative, measuring, magnitude  
...

Epoch: 5/5

Loss: 2.3859899044036865

often | more, are, common, or, some  
s | one, three, and, was, of  
they | them, are, would, be, but  
its | the, an, in, from, which  
who | his, he, had, by, himself  
eight | nine, seven, four, one, five  
no | info, import, jargon, duplicate, done  
over | the, years, and, three, four  
dr | starring, nine, producer, michael, scientist  
defense | armed, defence, nato, military, agency  
pope | papal, church, catholic, gregory, papacy  
mainly | parts, largest, region, largely, were

brother | son, daughter, wife, mother, consort  
account | accounts, critical, value, savings, debt  
engineering | electronics, technology, disciplines, engineers, technologies  
question | argument, whether, answers, how, reject  
...

Epoch: 5/5

Loss: 2.3489036560058594

when | to, would, but, was, him  
where | at, to, in, he, the  
an | a, s, to, of, was  
may | or, are, be, to, as  
during | after, in, was, war, s  
were | was, had, to, on, the  
would | when, did, able, him, had  
by | in, the, was, and, a  
shown | usually, if, called, show, closely  
universe | cosmology, marvel, galaxy, worlds, creator  
square | kilometers, adjacent, km, east, area  
writers | novelists, fiction, deaths, births, literary  
brother | son, mother, sister, father, younger  
ice | hockey, temperatures, snow, temperature, frozen  
placed | should, stick, others, would, usually  
shows | tv, episodes, they, game, movies  
...

Epoch: 5/5

Loss: 2.3585071563720703

more | some, than, as, both, often  
use | used, using, systems, these, uses  
about | this, have, up, that, zero  
at | the, in, on, s, of  
of | in, the, a, and, to  
and | the, to, of, a, in  
used | uses, use, and, other, a  
up | would, which, about, to, move  
mainly | include, largest, mostly, throughout, were  
mathematics | mathematical, mathematicians, algebra, calculus, algebraic  
arts | art, disciplines, martial, school, aikido  
institute | university, research, polytechnic, studies, museum  
grand | prix, at, masters, wins, duke  
paris | du, le, jean, france, des  
ocean | islands, pacific, atlantic, atolls, coral  
joseph | born, edward, b, james, chemist  
...

## Visualizing the word vectors

Below we'll use T-SNE to visualize how our high-dimensional word vectors cluster together. T-SNE is used to project these vectors into two dimensions while preserving local structure. Check out [this post from Christopher Olah \(http://colah.github.io/posts/2014-10-Visualizing-MNIST/\)](http://colah.github.io/posts/2014-10-Visualizing-MNIST/) to learn more about T-SNE and other ways to visualize high-dimensional data.

```
%matplotlib inline
%config InlineBackend.figure_format = 'retina'

import matplotlib.pyplot as plt
from sklearn.manifold import TSNE
```

```
# getting embeddings from the embedding layer of our model, by name
embeddings = model.in_embed.weight.to('cpu').data.numpy()
```

```
viz_words = 380
tsne = TSNE()
embed_tsne = tsne.fit_transform(embeddings[:viz_words, :])
```

```
fig, ax = plt.subplots(figsize=(16, 16))
for idx in range(viz_words):
    plt.scatter(*embed_tsne[idx, :], color='steelblue')
    plt.annotate(int to_vocab[idx], (embed_tsne[idx, 0], embed_tsne[idx, 1]), alpha
```

In [ ]: